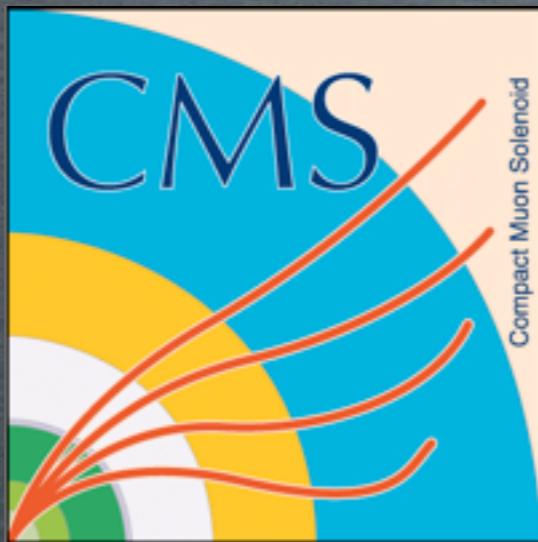


# LOBSTER: SCALING OPPORTUNISTIC CMS WORKFLOWS TO 10K CORES

ANNA WOODARD, MATTHIAS WOLF, CHARLES MUELLER, NIL  
VALLS, BEN TOVAR, PATRICK DONNELLY, PETER IVIE, PAUL  
BRENNER, DOUGLAS THAIN, KEVIN LANNON, MICHAEL HILDRETH

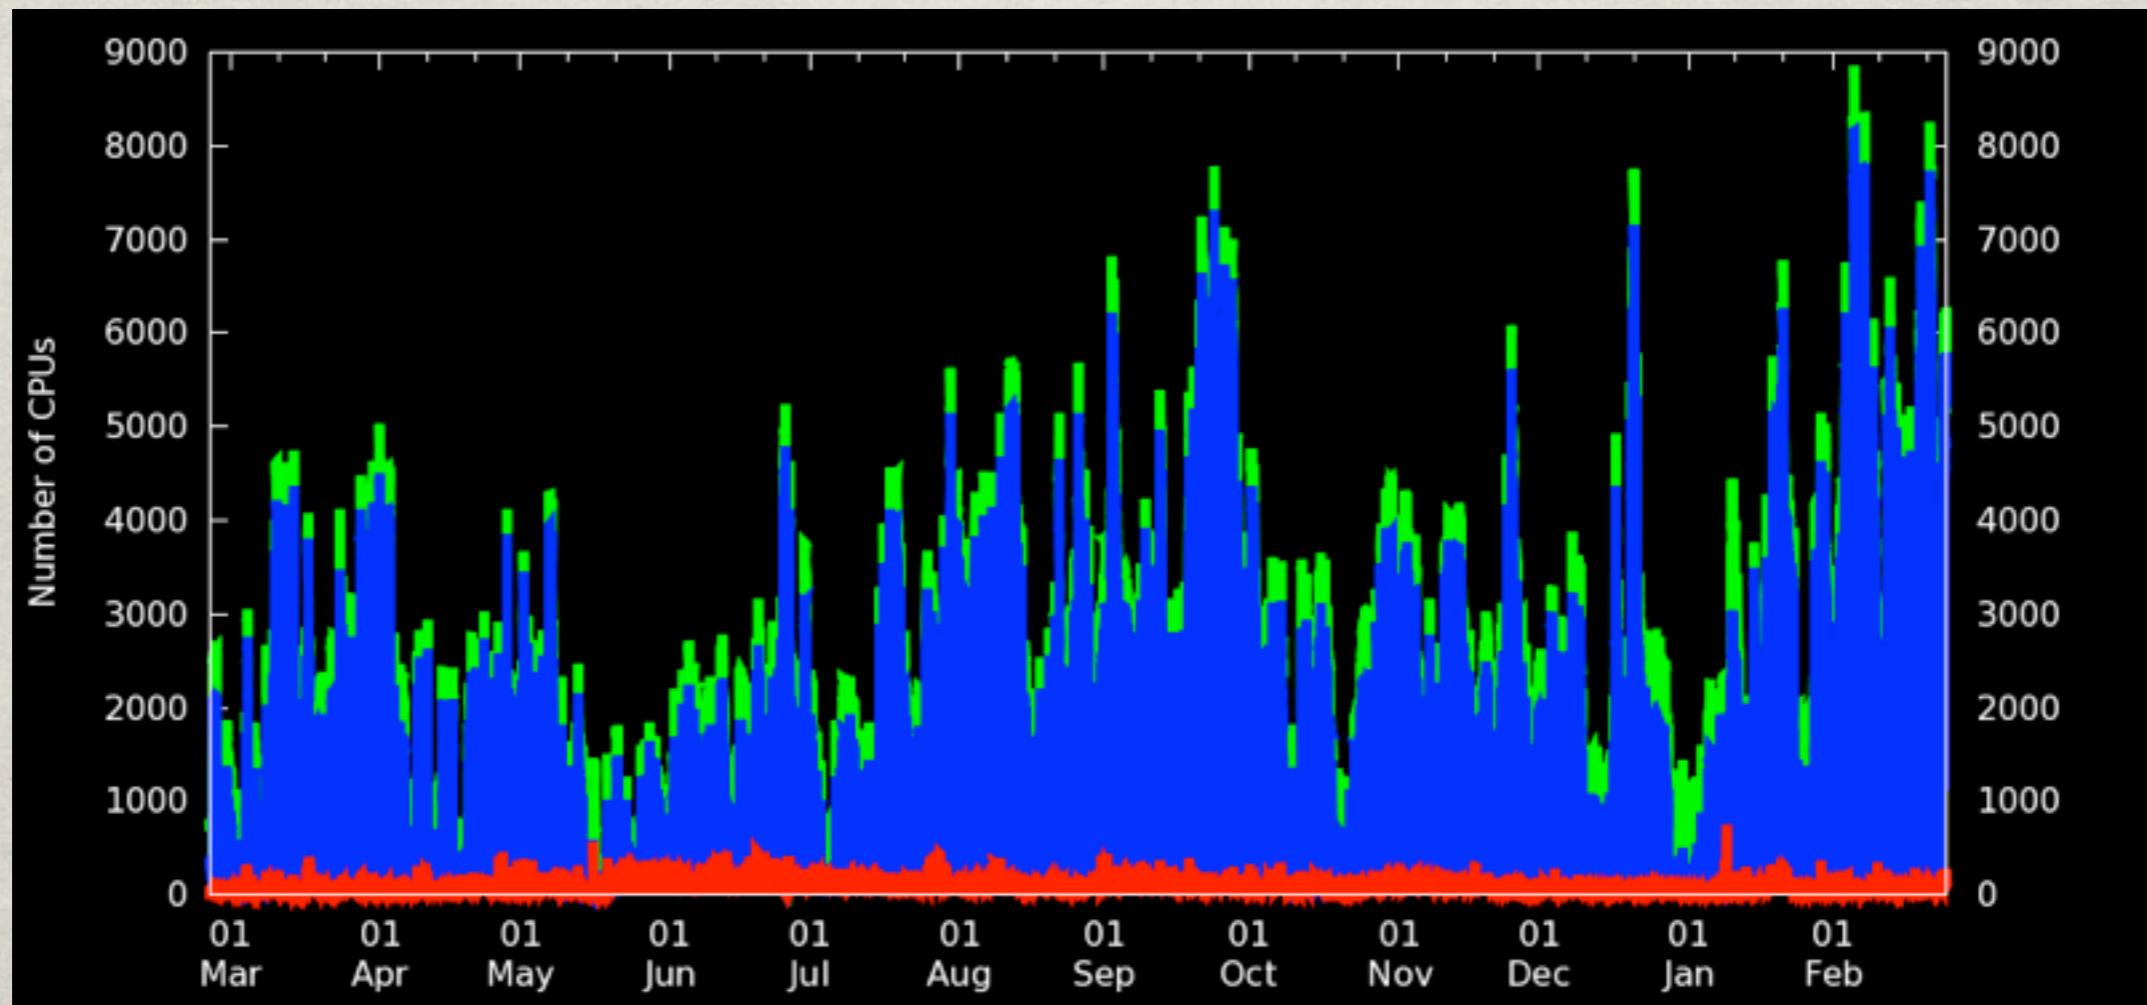


# WHAT IS LOBSTER?

- ⌘ Large-scale Opportunistic Batch Submission Toolkit for Exploiting Resources
- ⌘ Workflow submission and management tool written from scratch by two ND grad students (Matthias Wolf and Anna Woodard)
- ⌘ Borrows ideas from CRAB2/3 and grid-control
- ⌘ Based on CCTools suite (WorkQueue, Parrot, Chirp) from ND team led by Doug Thain
- ⌘ Primary goal: Get access to ND's opportunistic computing resources

# ND CRC RESOURCES

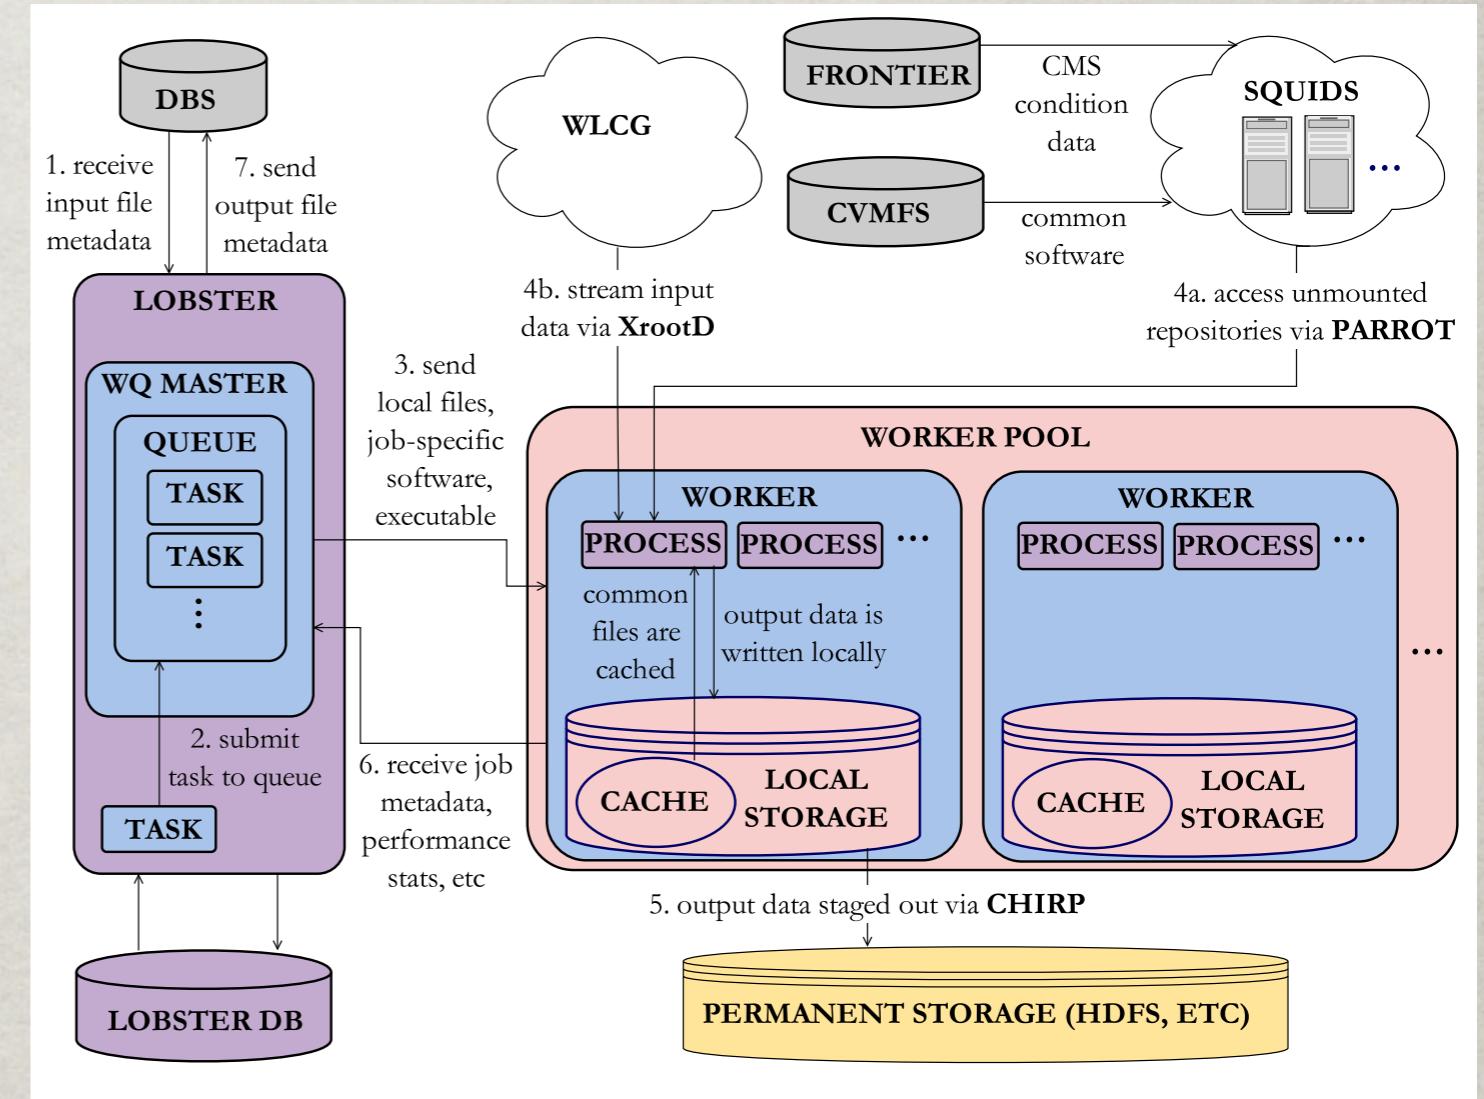
- ✿ ND Center for Research Computing houses ~21k CPU cores and 2.5 PB of storage
  - ✿ Most resources belong to individual PIs
  - ✿ Available for opportunistic usage when idle (evicted when owners reclaim resources)



# LOBSTER ARCHITECTURE

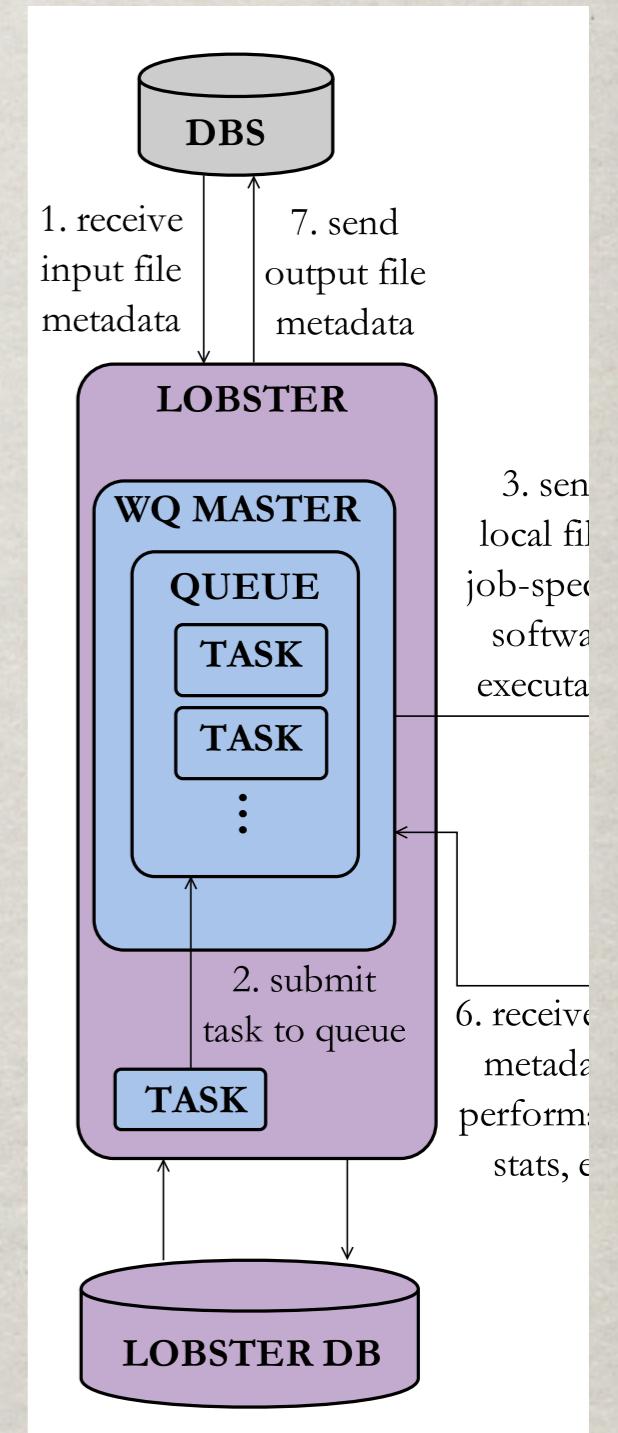
## >Main components

- Scheduling: schedules and dispatches jobs
- Data: managing input/output data and software
- Execution: runs tasks on opportunistic resources
- Master-worker architecture



# SCHEDULING

- ✿ Lobster master uses DBS to build database of work to be done
  - ✿ Work broken down into smallest sensible quantum: jobit
  - ✿ Lobster master schedules assigns jobits to tasks and schedules in Work Queue (WQ)
- ✿ WQ master handles distributing tasks to workers and tracks task progress
  - ✿ Optionally, system may include “foremen” to mediate between WQ master and workers
- ✿ Lobster communicates with WQ master to track jobit completion



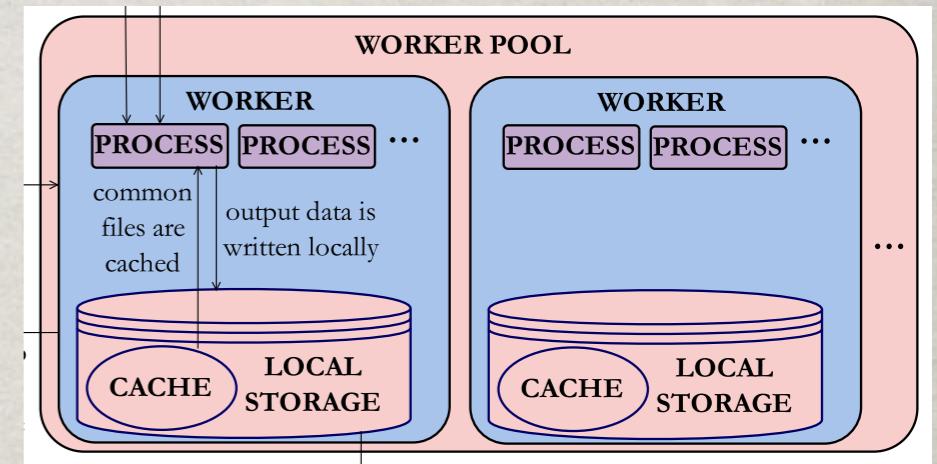
# DATA



- ✿ Leverage wide variety of tools (CVMFS, Parrot, Chirp, XrootD, WQ) to distribute data to jobs
  - ✿ CMSSW distributed to workers via CVMFS+Parrot (squid cache, worker cache)
  - ✿ Job scripts and sandbox transferred via WQ (worker cache)
  - ✿ Conditions via Frontier (squid cache)
  - ✿ Input data delivered via AAA (XrootD) or ND T3 storage (XrootD or Chirp)
  - ✿ Outputs transferred via Chirp

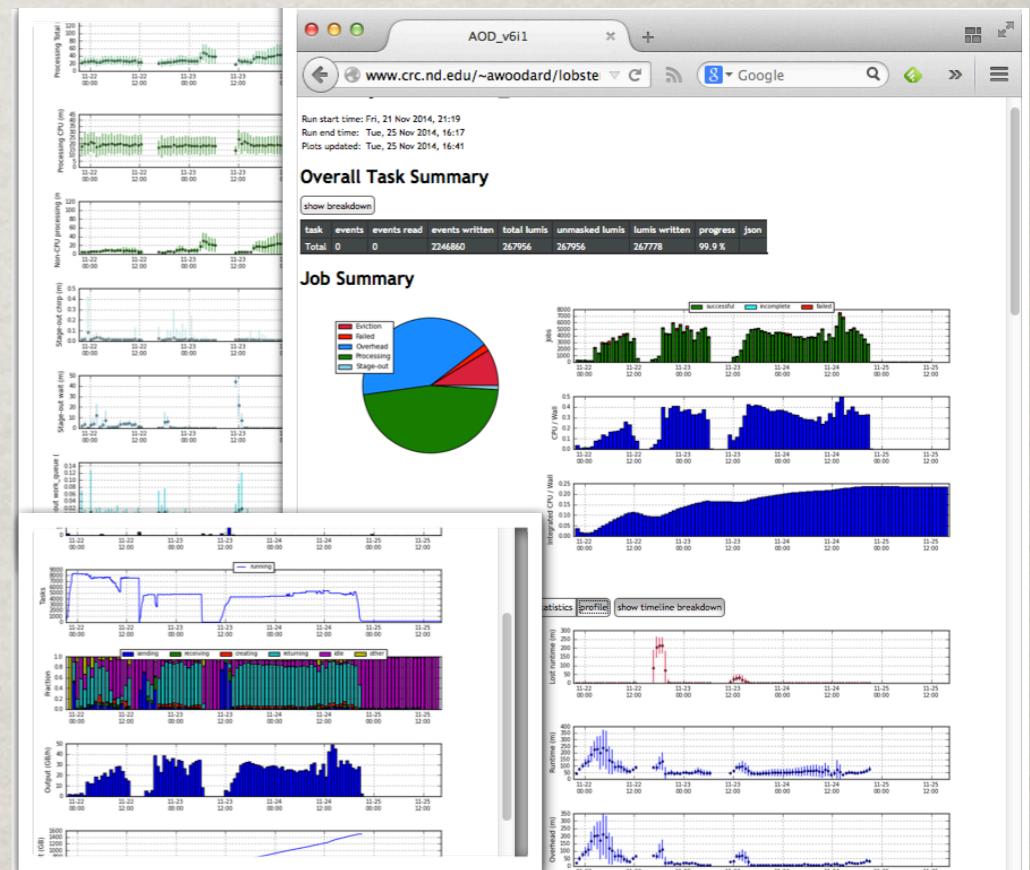
# EXECUTION: WORKERS

- ✿ Can be submitted via whatever batch system is available (HTCondor, SGE, PBS, etc.)
- ✿ CCTools includes tools for managing worker pools
- ✿ Responsible for configuring resource to accept CMS tasks (setup CMSSW, etc.)
- ✿ Holds resources and runs tasks for master until work is finished or worker is evicted
- ✿ Multicore workers will run multiple tasks in parallel, sharing local cache for CVMFS and WQ files



# MONITORING

- ✿ Opportunistic resources change dynamically (chaotically)
  - ✿ Resources come and go depending on owner activity
  - ✿ Heterogeneous quality
  - ✿ Can fail randomly
- ✿ Monitoring critical to Lobster success
  - ✿ Lobster tracks time stamps of every phase of task setup and execution
  - ✿ Collects information in plots and tables on webpage
  - ✿ Gives comprehensive picture of system components so that bottlenecks and failures can be diagnosed and overcome

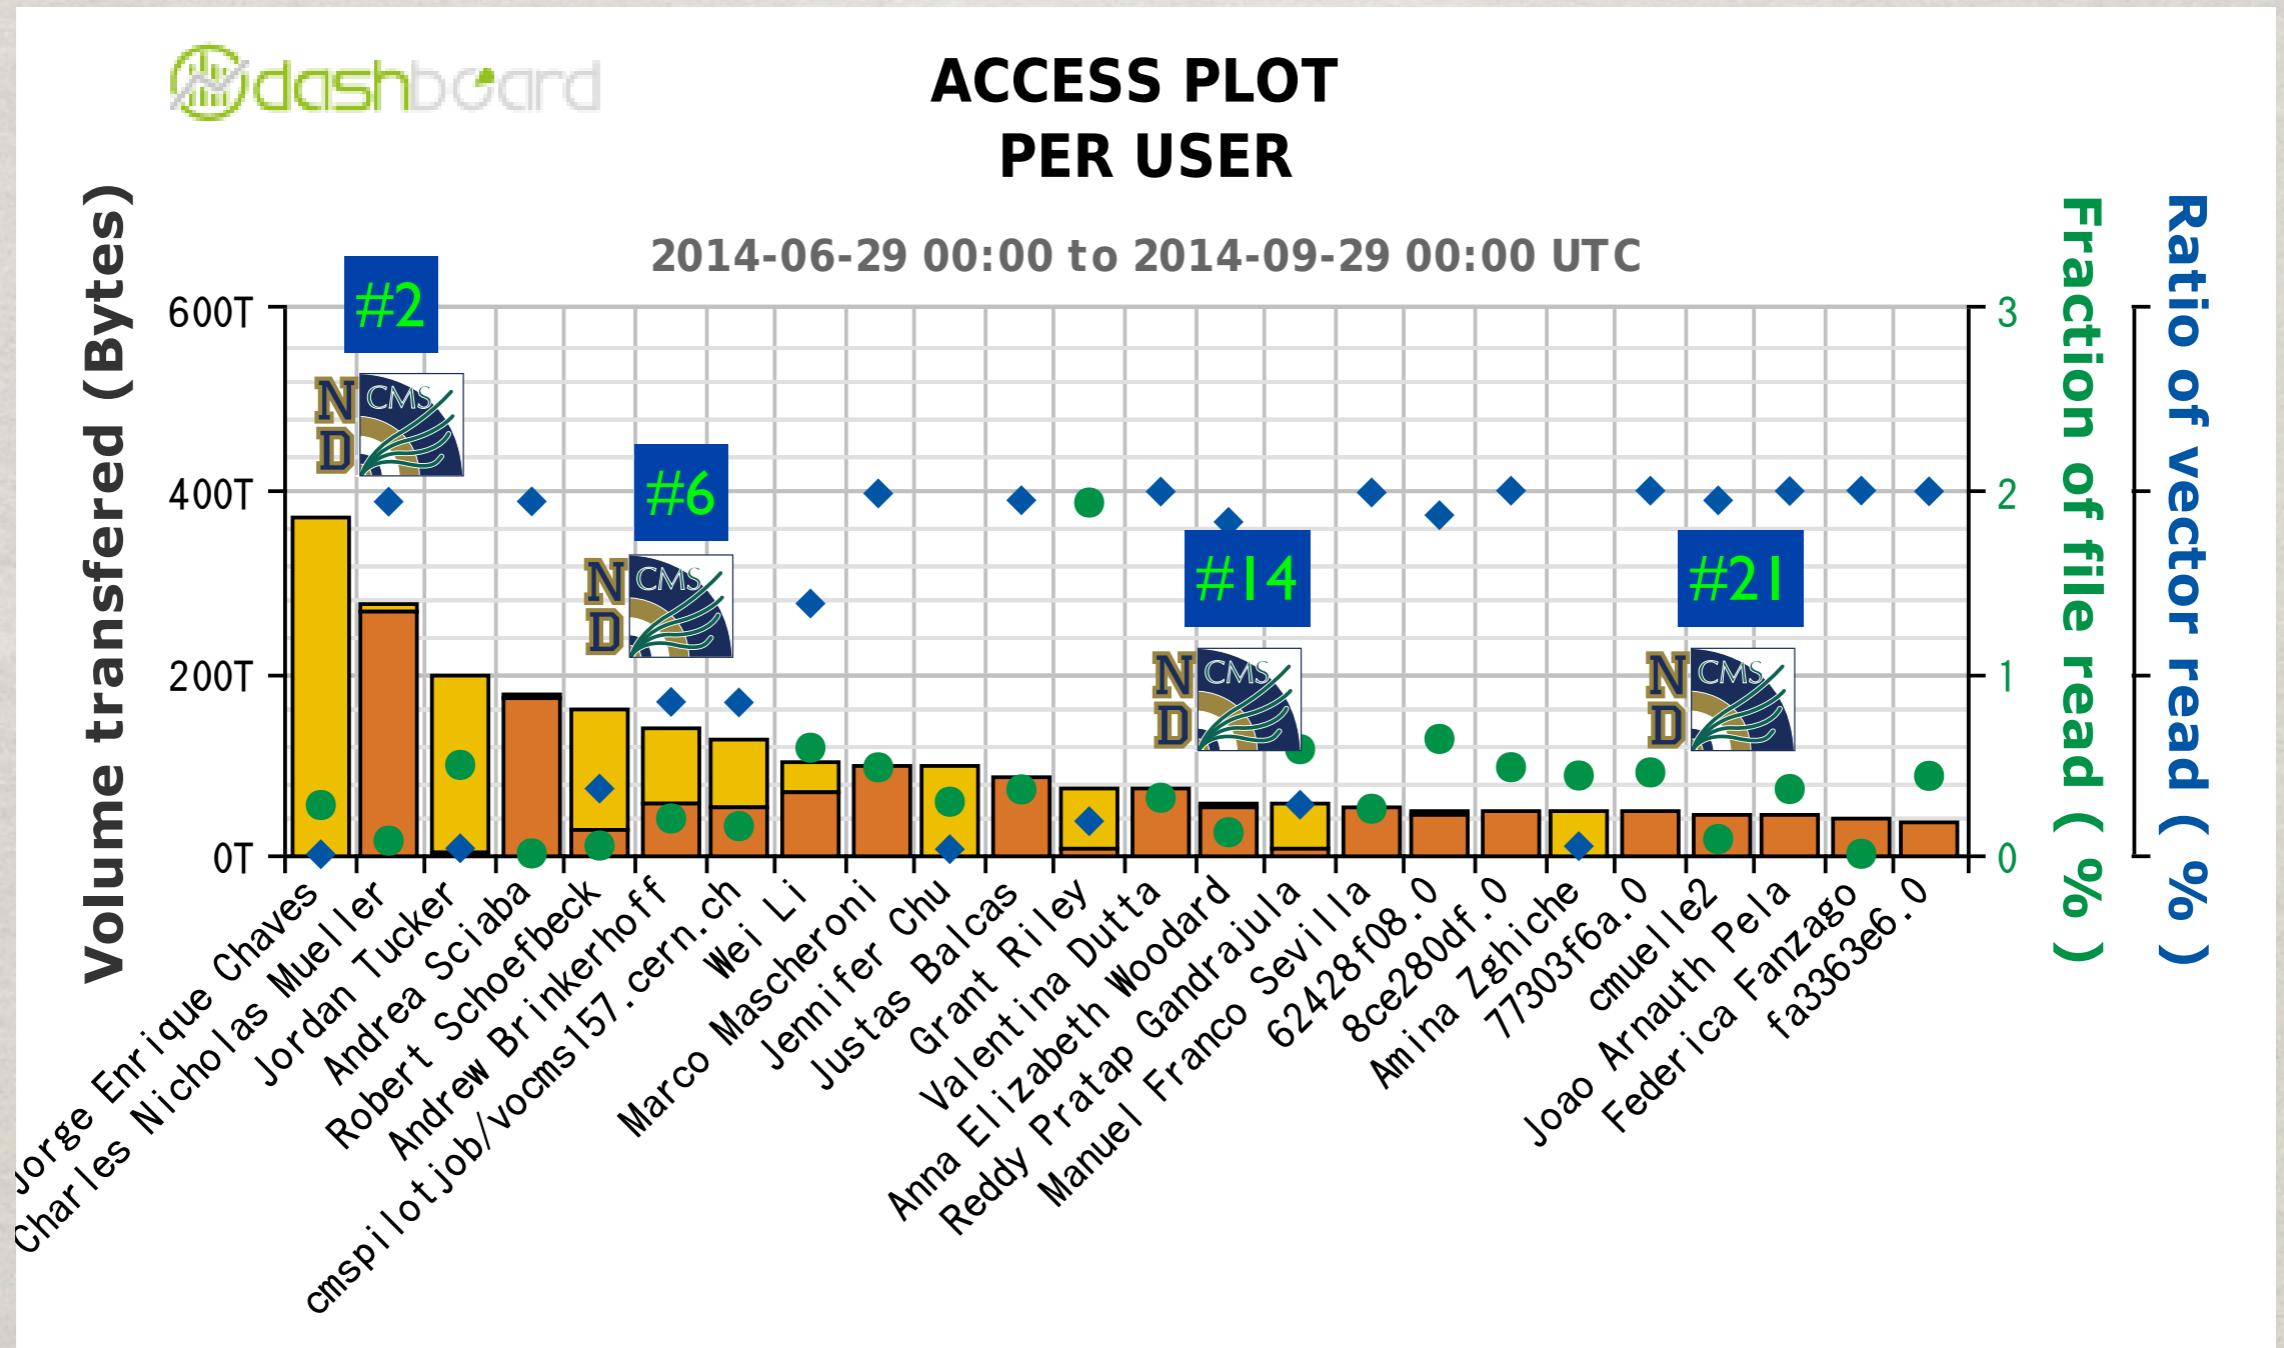


# DESIGN MOTIVATIONS

- ✿ Maximize opportunistic spirit!
- ✿ CCTools suite operates completely in user space: no admin intervention required to use opportunistic resources
- ✿ Eviction requires agility: Decouple job size from output size and user task management
  - ✿ Lobster works in jobits, tracks splitting, handles resubmission without user intervention
  - ✿ Often leads to really small output files--merged automatically
  - ✿ When processing multiple datasets, jobit execution randomized to level load on AAA
- ✿ Persistence pays off
  - ✿ Workers try very hard to get tasks started: Use local CVMFS if present, switch to Parrot if not, etc.
  - ✿ Lobster retries failed tasks until you stop it: needs to be resilient against transient failure in opportunistic system
  - ✿ Work Queue Pool: if workers die/lost, resubmit more

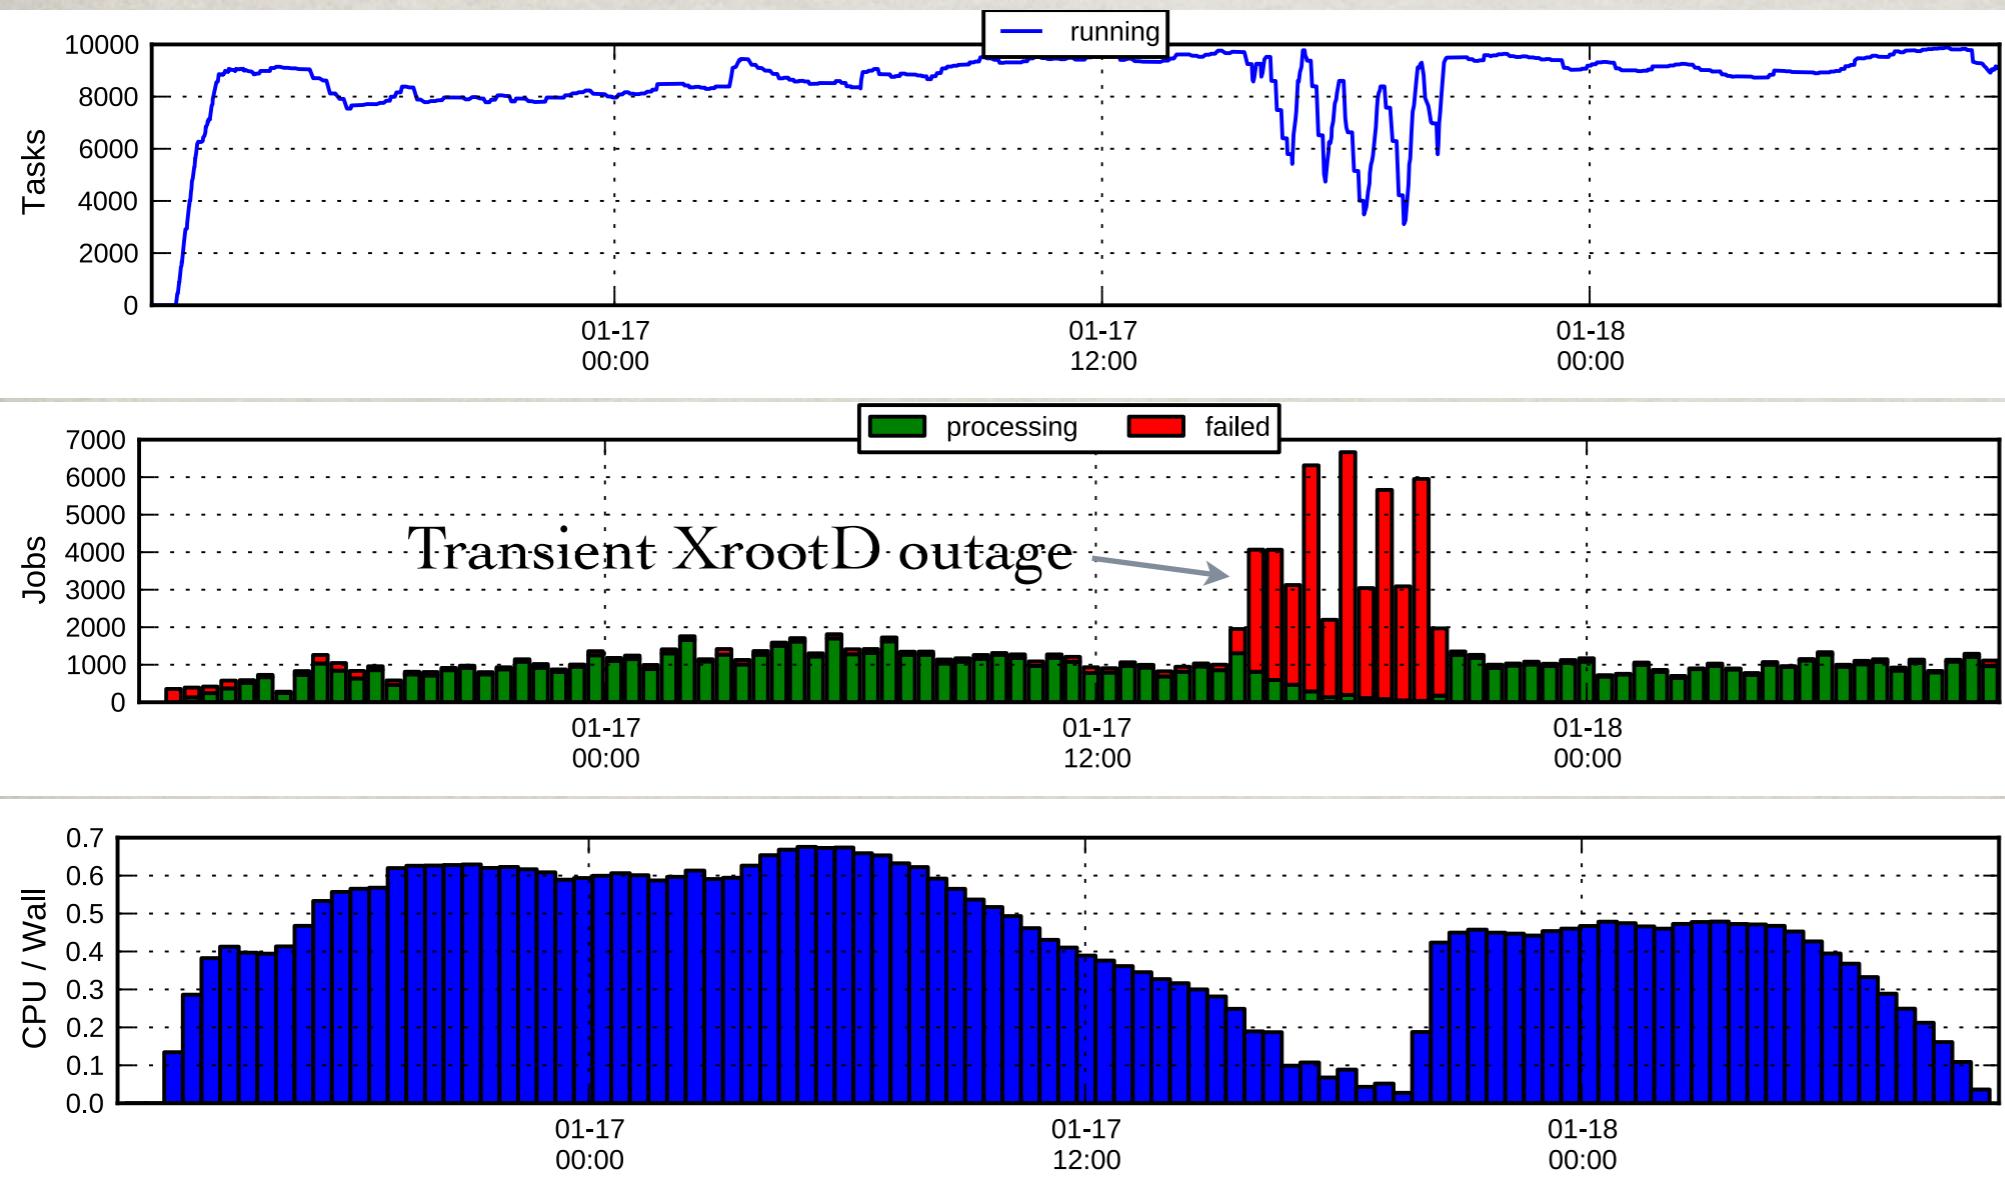
# PROCESSING LOTS OF DATA

- ND users running only on ND resources competitive with CSA14 activity



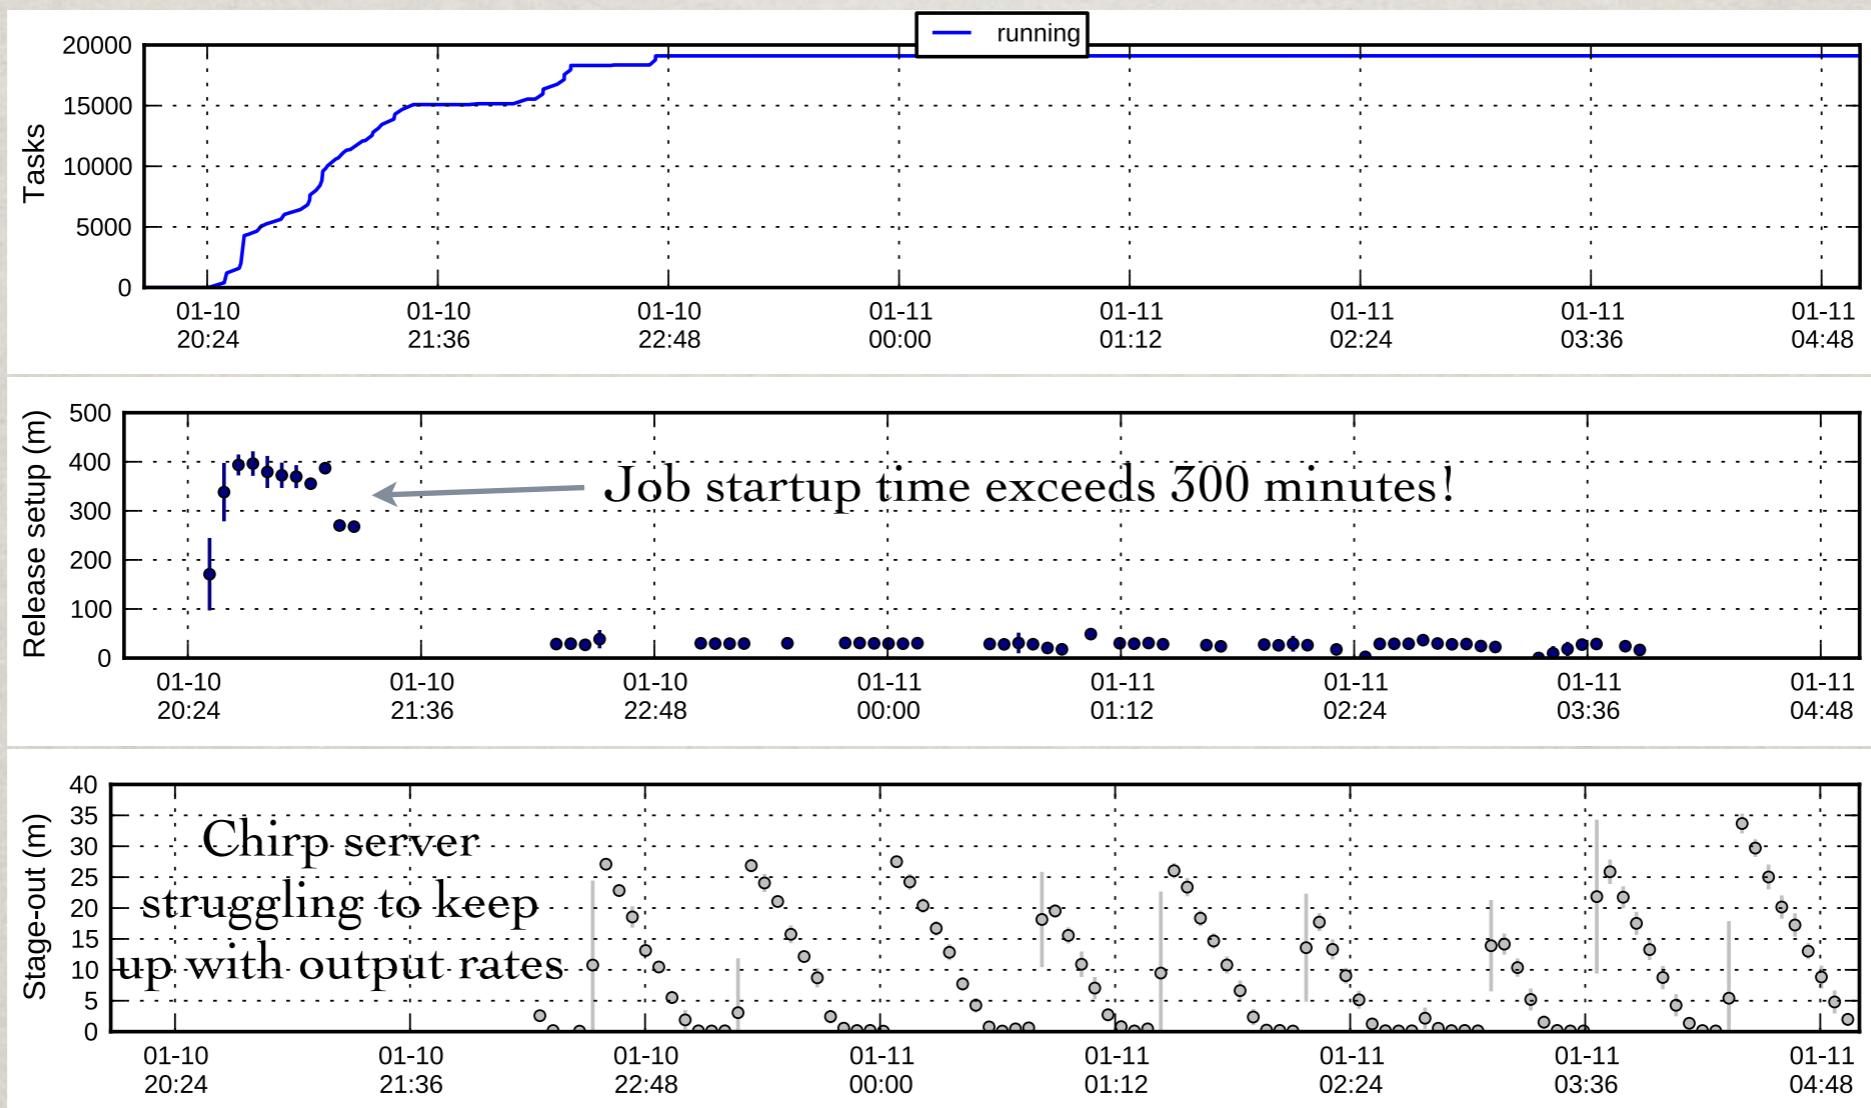
# REACHING ~10K RUNNING JOBS

- This is ~10% of size of CMS global pool
- Comparable to scale of one US CMS T2 site
- More than total of all US CMS dedicated T3 resources



# CHALLENGES

- ✿ Bottlenecks when running at large scales
  - ✿ Job overhead time increases non-linearly: bottleneck in squid cache?
  - ✿ Output Chirp server can get overloaded and fail



# NEAR TERM FUTURE

- ✿ Resolve bottlenecks preventing scaling to 20k cores
  - ✿ Also important for decreasing overhead to minimize eviction losses
- ✿ Explore possibilities for Lobster to dynamically adapt to running conditions
  - ✿ E.g. Automatically adjust jobits/task to optimize for current running conditions (eviction)
- ✿ Improvements in reliability and robustness
  - ✿ E.g. Run Chirp server as service instead of user process to keep users from overwhelming login node with Chirp processes

# CONCLUSIONS

- ✿ Lobster has enabled ND team to exploit opportunistic campus resources to 10k core scale
- ✿ Successful collaboration between physics and CS teams
  - ✿ Learned a number of useful things
  - ✿ Motivated improvements to CCTools suite
- ✿ Potential for lessons learned and innovations to be translated more widely to CMS
  - ✿ Anna and Matthias now working as CRAB3 developers
  - ✿ See Brian's talk for longer term vision